

Counting Bees using Mask R-CNN

Andreas Wittmann

2018/09/07 (updated: 2018-09-06)


Motivation

- My wife is a beekeeper
- She would like to know how high the bee traffic is
- Can we count the number of bees coming and going over time?
- **Goals:**
 - Analyze bee traffic over time using data collected from a webcam
 - Develop an app which shows the number of bees on a real-time video

Sounds like an application for deep learning...

CNN

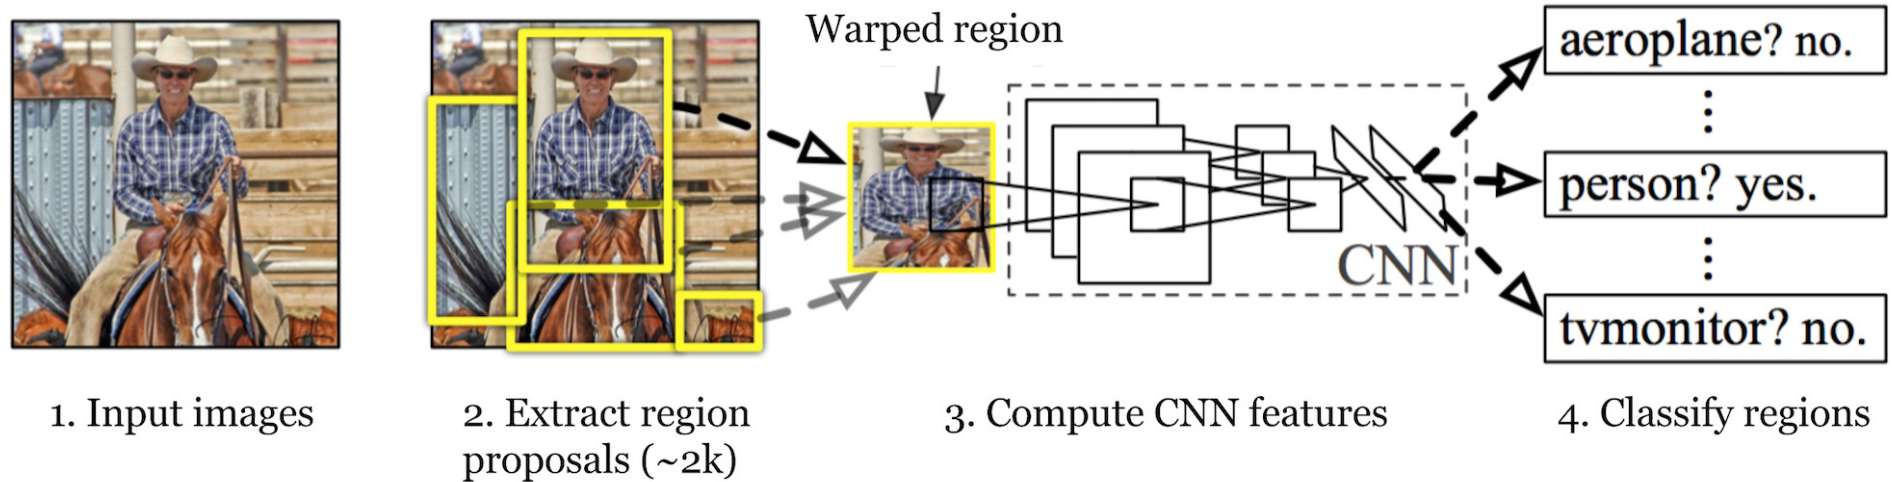
- **Convolutional Neural Networks** (CNNs): gold standard for image classification
- Usually there is an image with one single object instance and the focus of the task is to say what that image is showing
- We need **object detection/segmentation** here
- For counting objects we have the following challenges:
 - type of the objects to be counted
 - overlapping
 - perspective view
 - the minimum size of detected objects
 - training and testing speed



R-CNN

- **Region-based Convolutional Neural Networks (R-CNN)**
- Girshick et al., 2014
- Purpose is to solve the problem of object detection
- Given a certain image, we want to be able to draw bounding boxes over all of the objects
- **Region proposal step:**
 - *Selective search* generates 2000 different object regions with the highest probability of containing an object
 - Proposals are “warped” into an image size that can be fed into a pre-trained CNN (e.g. AlexNet)
 - A feature vector is extracted for each region
- **Classification step:**
 - Use a trained SVM to classify each feature vector for each class independently
 - Use a simple bounding-box regression with each feature vector to obtain the most accurate coordinates

R-CNN



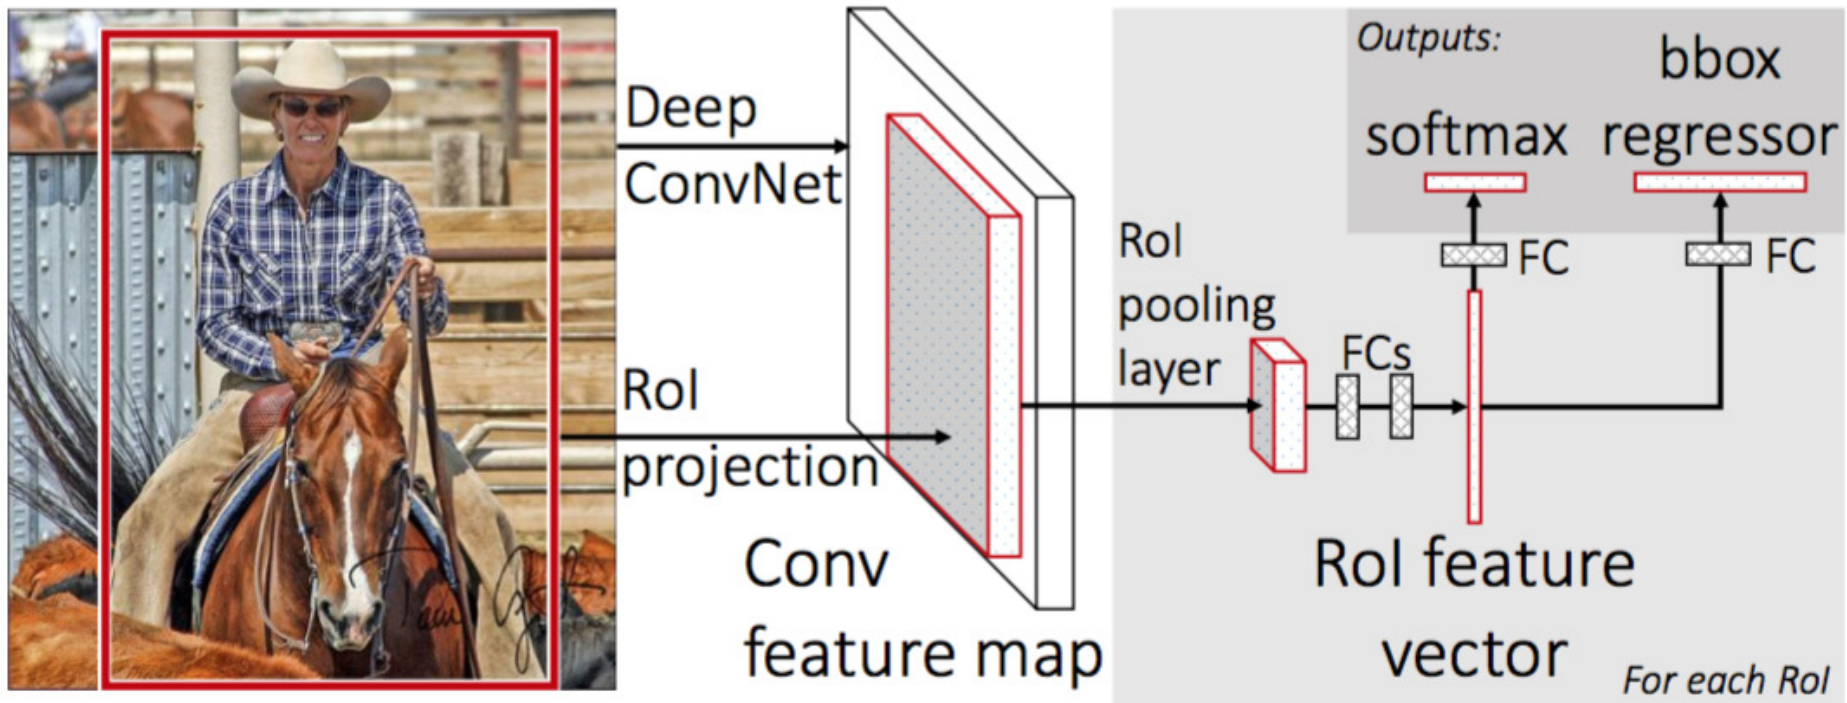
Speed Bottleneck:

- Running selective search to propose 2000 region candidates for every image
- Generating the CNN feature vector for every image region ($N \text{ images} * 2000$)

Fast R-CNN

- **Girshick et al., 2015**: Union three independent models into one, named *Fast R-CNN*
- That Last layer of a pre-trained CNN is replaced by a RoI pooling Layer
- **RoI (Region of Interest) Pooling:**
 - Run CNN once per image (instead of 2000 times) and share computation across proposals
 - Many region proposals of the same images are highly overlapped
- **Combine All Models into One Network:**
 - Jointly train the CNN, classifier, and bounding box regressor in a single model.

Fast R-CNN



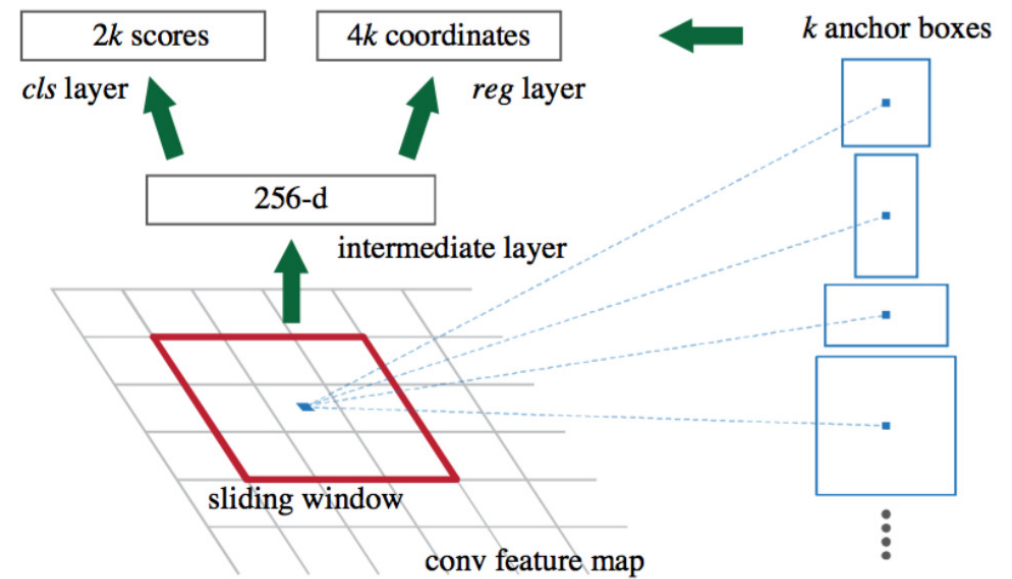
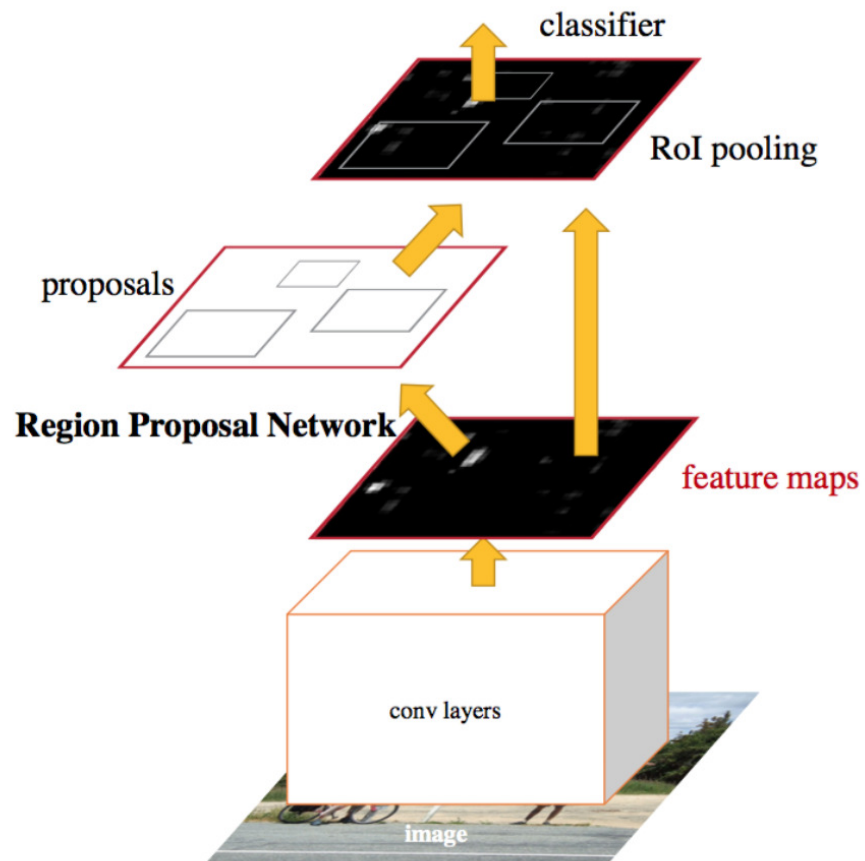
Speed Bottleneck:

- Fast R-CNN is faster in both training and testing time.
- However, the improvement is not dramatic because the region proposals are generated separately by another model and that is very expensive.

Faster R-CNN

- Integrate the region proposal algorithm into the CNN model for speedup
- [Ren et al., 2016](#): construct a single model composed of
 - **Region proposal network (RPN)**: set of object proposals with objectness score
 - fast R-CNN with shared convolutional feature layers:
 - Slide a small (3 x 3) network over the convolutional feature map
 - Map to lower-dimensional feature 256-d
 - Box-regression layer (reg)
 - Box-classification layer (cls)
 - **Anchors**: Predict k proposals relative to k reference boxes

Faster R-CNN



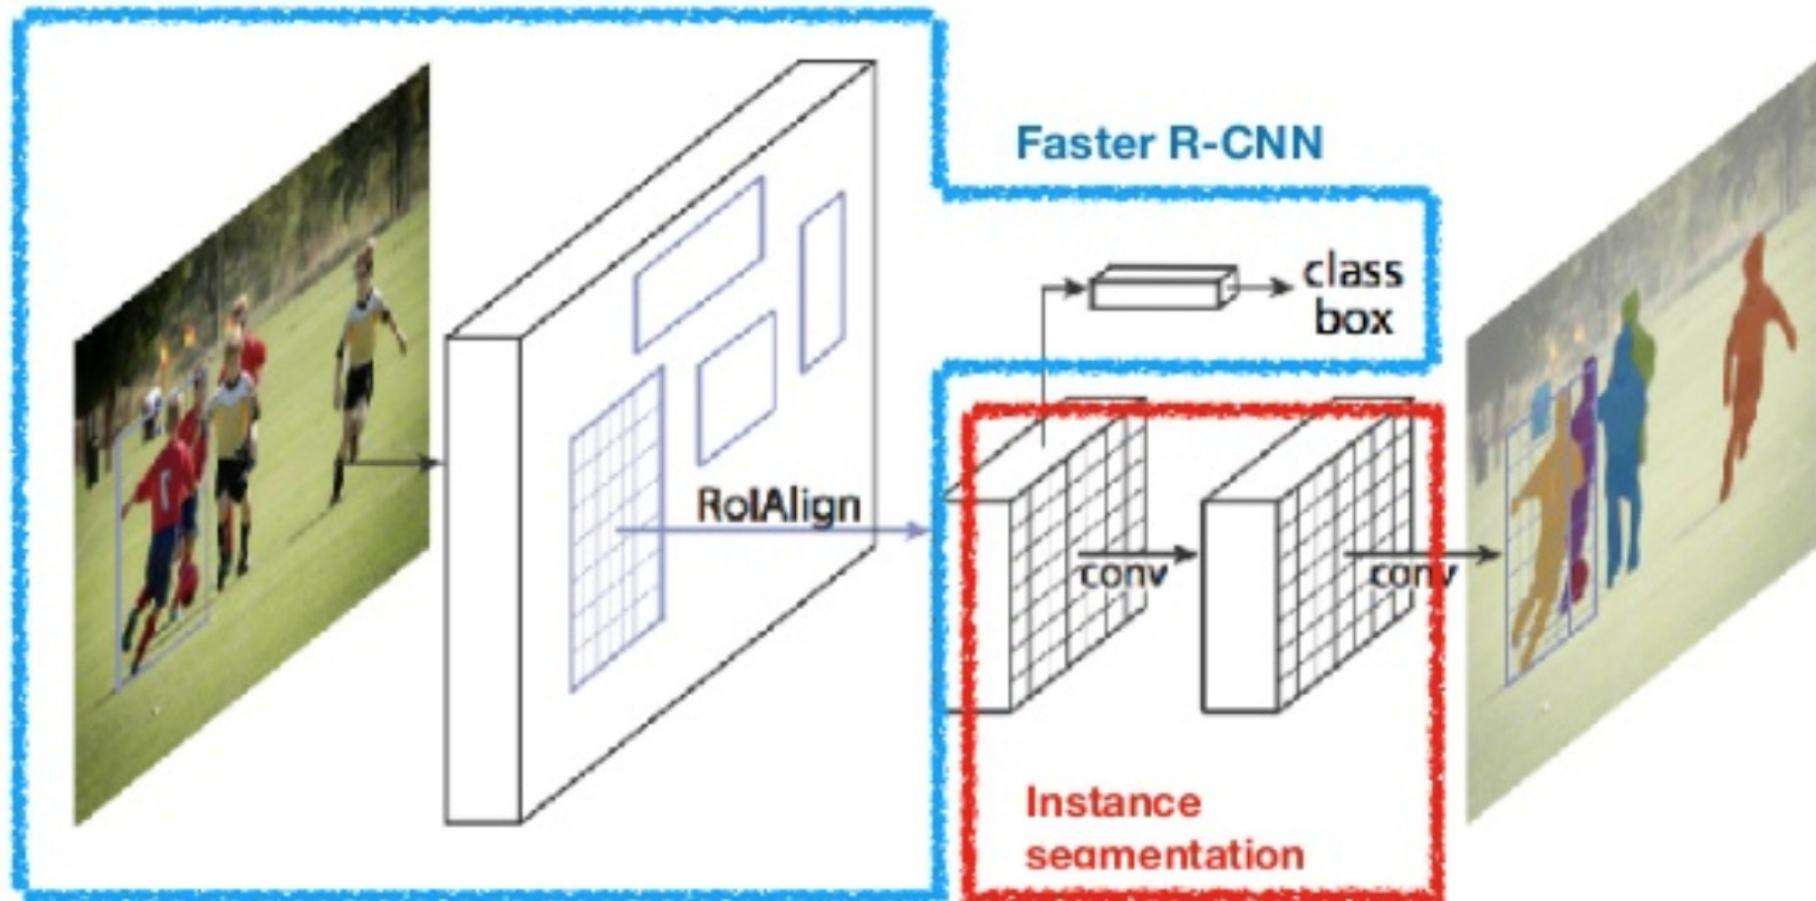
Mask R-CNN

- So far: CNN features to locate different objects in an image with bounding boxes.
- Mask R-CNN [He et al., 2017](#) extends Faster R-CNN to pixel-level image segmentation.
- Use **RoIAlign** instead of RoI Pooling: avoids pixel-wise rounding problems
- Output: **binary mask** that says whether or not a given pixel is part of an object
- It can also color pixels in the bounding box that correspond to that class

Mask R-CNN


Combination of

- Faster R-CNN for object detection (class and bounding box)
- FCN (Fully Convolutional Network): pixel wise boundary




Building a Mask R-CNN Model for Counting Bees

- A Mask R-CNN implementation on Python 3, Keras and Tensorflow is used
- 4 videos next to the hive in different perspectives with a length of about 1 minute where recorded
- From each video 20 frames where extracted
- These are devided into train (16 images), validation (8 images) and test (56 images) dataset
- Annotating the train data using an online tool



Anchors



Training the model

- COCO Dataset consists of 330,000 images with more than 200,000 images labeled with 80 object categories
- The COCO pre-trained model is used as the checkpoint for transfer learning
- **Transfer learning:** adaptation of pretrained models to similar or moderately different tasks, by finetuning parameters of the pretrained models
- Training the model on a server with 250 gigabyte RAM and 32 cores (without GPU) for 10 epochs took about 12 hours

Results (1)



Results (2)



Summary

- As a proof of concept the approach seems to work
- But...
 - We need to label more data to improve the results
 - The prediction seems somehow slow
 - Maybe use polygons instead of single pixels to get it work for near perspectives
 - Try other approaches like YOLO or SSD as they can achieve higher framerates

Literature

- <https://blog.athelas.com/a-brief-history-of-cnns-in-image-segmentation-from-r-cnn-to-mask-r-cnn-34ea83205de4>
- https://medium.com/@umerfarooq_26378/from-r-cnn-to-mask-r-cnn-d6367b196cfd
- http://matpalm.com/blog/counting_bees/
- <https://softwaremill.com/counting-objects-with-faster-rcnn/>
- <https://lilianweng.github.io/lil-log/2017/12/31/object-recognition-for-dummies-part-3.html>
- R-CNN: <https://arxiv.org/abs/1311.2524>
- Fast R-CNN: <https://arxiv.org/abs/1504.08083>
- Faster R-CNN: <https://arxiv.org/abs/1506.01497>
- Mask R-CNN: <https://arxiv.org/abs/1703.06870>
- Mask R-CNN Implementation: https://github.com/matterport/Mask_RCNN